

# Computer Network Assignment 3

Yiping Deng

May 10, 2019

## Problem 3

a)

Using topology defined in Listing 2, we can execute the ping and traceroute in mininet and verify the topology:

Listing 1: Mininet output for a)

```
1 mininet> h1 ping6 2001:638:709:b::1
2 PING 2001:638:709:b::1(2001:638:709:b::1) 56 data bytes
3 64 bytes from 2001:638:709:b::1: icmp_seq=1 ttl=62 time=0.125 ms
4 64 bytes from 2001:638:709:b::1: icmp_seq=2 ttl=62 time=0.060 ms
5 64 bytes from 2001:638:709:b::1: icmp_seq=3 ttl=62 time=0.060 ms
6 ^C
7 --- 2001:638:709:b::1 ping statistics ---
8 3 packets transmitted, 3 received, 0% packet loss, time 2050ms
9 rtt min/avg/max/mdev = 0.060/0.081/0.125/0.032 ms
10
11
12 mininet> h2 ping6 2001:638:709:a::1
13 PING 2001:638:709:a::1(2001:638:709:a::1) 56 data bytes
14 64 bytes from 2001:638:709:a::1: icmp_seq=1 ttl=62 time=0.058 ms
15 64 bytes from 2001:638:709:a::1: icmp_seq=2 ttl=62 time=0.062 ms
16 64 bytes from 2001:638:709:a::1: icmp_seq=3 ttl=62 time=0.073 ms
17 64 bytes from 2001:638:709:a::1: icmp_seq=4 ttl=62 time=0.064 ms
18 64 bytes from 2001:638:709:a::1: icmp_seq=5 ttl=62 time=0.060 ms
19 ^C
20 --- 2001:638:709:a::1 ping statistics ---
21 5 packets transmitted, 5 received, 0% packet loss, time 4087ms
22 rtt min/avg/max/mdev = 0.058/0.063/0.073/0.008 ms
23
24
25
26 mininet> h1 traceroute 2001:638:709:b::1
27 traceroute to 2001:638:709:b::1 (2001:638:709:b::1), 30 hops max, 80 byte packets
28 1 2001:638:709:a::f (2001:638:709:a::f) 0.033 ms 0.018 ms 0.008 ms
29 2 2001:638:709:f::2 (2001:638:709:f::2) 0.018 ms 0.009 ms 0.009 ms
30 3 2001:638:709:b::1 (2001:638:709:b::1) 0.019 ms 0.010 ms 0.010 ms
31
32
33
34 mininet> h2 traceroute 2001:638:709:a::1
35 traceroute to 2001:638:709:a::1 (2001:638:709:a::1), 30 hops max, 80 byte packets
36 1 2001:638:709:b::f (2001:638:709:b::f) 0.030 ms 0.007 ms 0.007 ms
37 2 2001:638:709:f::1 (2001:638:709:f::1) 0.019 ms 0.009 ms 0.008 ms
38 3 2001:638:709:a::1 (2001:638:709:a::1) 0.019 ms 0.010 ms 0.010 ms
```

Listing 2: Basic topology network

---

```

1  #!/usr/bin/env python
2  from mininet.cli import CLI
3  from mininet.net import Mininet
4  from mininet.nodelib import LinuxBridge
5  from mininet.log import setLogLevel
6
7  def set_ipv6(host, interface, ip):
8      return host.cmd("ip -6 addr add " + ip + " dev " + interface)
9
10 def set_ipv6_gateway(host, gateway):
11     return host.cmd("ip -6 route add default via " + gateway)
12
13 def set_ipv6_device_route(host, subnet, interface):
14     return host.cmd("ip -6 route add " + subnet + " dev " + interface)
15
16 def set_ipv6_gateway_route(host, subnet, gateway, interface):
17     # useful for the forward table
18     return host.cmd("ip -6 route add " + subnet +
19                     " via " + gateway + " dev " + interface)
20
21
22
23
24 if __name__ == '__main__':
25     setLogLevel('info')
26
27     net = Mininet(switch=LinuxBridge, controller=None)
28
29     h1 = net.addHost('h1', ip=None)
30
31     # create rest of the topology here
32     s1 = net.addSwitch('s1')
33     r1 = net.addHost('r1', ip = None)
34     s0 = net.addSwitch('s0')
35     r2 = net.addHost('r2', ip = None)
36     s2 = net.addSwitch('s2')
37     h2 = net.addHost('h2', ip = None)
38
39     # addr info
40     h1_eth0_addr = "2001:638:709:a::1/64"
41     r1_eth0_addr = "2001:638:709:a::f/64"
42     r1_eth1_addr = "2001:638:709:f::1/64"
43     r2_eth0_addr = "2001:638:709:f::2/64"
44     r2_eth1_addr = "2001:638:709:b::f/64"
45     h2_eth0_addr = "2001:638:709:b::1/64"
46
47     # subnet info
48     subnet1_addr = "2001:638:709:a::/64"
49     subnet2_addr = "2001:638:709:f::/64"
50     subnet3_addr = "2001:638:709:b::/64"
51
52     # link the router and switches
53     net.addLink(h1, s1)
54     net.addLink(s1, r1)

```

```

55     net.addLink(r1, s0)
56     net.addLink(s0, r2)
57     net.addLink(r2, s2)
58     net.addLink(s2, h2)
59
60     # configure IPv6 addresses and forwarding table entries here
61
62     # set ipv6 addr and gateway
63     print "setup ipv6"
64     set_ipv6(h1, "h1-eth0", h1_eth0_addr)
65     set_ipv6(r1, "r1-eth0", r1_eth0_addr)
66     set_ipv6(r1, "r1-eth1", r1_eth1_addr)
67     set_ipv6(r2, "r2-eth0", r2_eth0_addr)
68     set_ipv6(r2, "r2-eth1", r2_eth1_addr)
69     set_ipv6(h2, "h2-eth0", h2_eth0_addr)
70     print "done setup ipv6"
71
72     # setup the local route
73     print "setup forward rules"
74     set_ipv6_device_route(h1, subnet1_addr, "h1-eth0")
75     set_ipv6_device_route(r1, subnet1_addr, "r1-eth0")
76     set_ipv6_device_route(r1, subnet2_addr, "r1-eth1")
77     set_ipv6_device_route(r2, subnet2_addr, "r2-eth0")
78     set_ipv6_device_route(r2, subnet3_addr, "r2-eth1")
79     set_ipv6_device_route(h2, subnet3_addr, "h2-eth0")
80     print "done setup forward rules"
81
82     # setup the gateway
83     print "setup gateway"
84     set_ipv6_gateway(h1, r1_eth0_addr[:-3])
85     set_ipv6_gateway(h2, r2_eth1_addr[:-3])
86     print "done setup gateway"
87
88     # setup router forward table
89     print "setup router forward table"
90     set_ipv6_gateway_route(r1, subnet3_addr, r2_eth0_addr[:-3], "r1-eth1")
91     set_ipv6_gateway_route(r2, subnet1_addr, r1_eth1_addr[:-3], "r2-eth0")
92     print "done setup router forward table"
93
94     # router forwarding
95     print "enable forwarding for routers"
96     r1.cmd("sysctl -w net.ipv6.conf.all.forwarding=1")
97     r2.cmd("sysctl -w net.ipv6.conf.all.forwarding=1")
98     print "done enabling forwarding for routers"
99
100    print h1.cmd("ip -V")
101
102
103    net.start()
104    CLI(net)
105    net.stop()

```

b)

Using the modified Listing 4, we can use the traceroute to show the asymmetric path:

Listing 3: Traceroute of b)

---

```

1 mininet> h1 traceroute 2001:638:709:b::1
2 traceroute to 2001:638:709:b::1 (2001:638:709:b::1), 30 hops max, 80 byte packets
3  1  2001:638:709:a::f (2001:638:709:a::f)  0.042 ms  0.009 ms  0.006 ms
4  2  2001:638:709:f::2 (2001:638:709:f::2)  0.052 ms  0.011 ms  0.008 ms
5  3  2001:638:709:b::1 (2001:638:709:b::1)  0.117 ms  0.013 ms  0.010 ms
6
7
8 mininet> h2 traceroute 2001:638:709:a::1
9 traceroute to 2001:638:709:a::1 (2001:638:709:a::1), 30 hops max, 80 byte packets
10 1  2001:638:709:b::e (2001:638:709:b::e)  0.033 ms  0.008 ms  0.007 ms
11 2  2001:638:709:e::1 (2001:638:709:e::1)  0.054 ms  0.011 ms  0.009 ms
12 3  2001:638:709:a::1 (2001:638:709:a::1)  0.025 ms  0.010 ms  0.010 ms

```

---

Listing 4: Asymmetric route

---

```

1  #!/usr/bin/env python
2  from mininet.cli import CLI
3  from mininet.net import Mininet
4  from mininet.nodelib import LinuxBridge
5  from mininet.log import setLogLevel
6
7  def set_ipv6(host, interface, ip):
8      return host.cmd("ip -6 addr add " + ip + " dev " + interface)
9
10 def set_ipv6_gateway(host, gateway):
11     return host.cmd("ip -6 route add default via " + gateway)
12
13 def set_ipv6_device_route(host, subnet, interface):
14     return host.cmd("ip -6 route add " + subnet + " dev " + interface)
15
16 def set_ipv6_gateway_route(host, subnet, gateway, interface):
17     # useful for the forward table
18     return host.cmd("ip -6 route add " + subnet +
19                     " via " + gateway + " dev " + interface)
20
21
22
23
24 if __name__ == '__main__':
25     setLogLevel('info')
26
27     net = Mininet(switch=LinuxBridge, controller=None)
28
29     h1 = net.addHost('h1', ip=None)
30
31     # create rest of the topology here
32     s1 = net.addSwitch('s1')
33     r1 = net.addHost('r1', ip = None)
34     s0 = net.addSwitch('s0')
35     r2 = net.addHost('r2', ip = None)
36     s2 = net.addSwitch('s2')
37     h2 = net.addHost('h2', ip = None)
38     r3 = net.addHost('r3', ip = None)
39     s3 = net.addSwitch('s3')
40     r4 = net.addHost('r4', ip = None)
41
42     # addr info

```

---

```

43     h1_eth0_addr = "2001:638:709:a::1/64"
44     r1_eth0_addr = "2001:638:709:a::f/64"
45     r1_eth1_addr = "2001:638:709:f::1/64"
46     r2_eth0_addr = "2001:638:709:f::2/64"
47     r2_eth1_addr = "2001:638:709:b::f/64"
48     h2_eth0_addr = "2001:638:709:b::1/64"
49     r3_eth0_addr = "2001:638:709:a::e/64"
50     r3_eth1_addr = "2001:638:709:e::1/64"
51     r4_eth0_addr = "2001:638:709:e::2/64"
52     r4_eth1_addr = "2001:638:709:b::e/64"
53
54     # subnet info
55     subnet1_addr = "2001:638:709:a::/64"
56     subnet2_addr = "2001:638:709:f::/64"
57     subnet3_addr = "2001:638:709:b::/64"
58     subnet4_addr = " 2001:638:709:e::/64"
59
60     # link the router and switches
61     net.addLink(h1, s1)
62     net.addLink(s1, r1)
63     net.addLink(r1, s0)
64     net.addLink(s0, r2)
65     net.addLink(r2, s2)
66     net.addLink(s2, h2)
67     net.addLink(s1, r3)
68     net.addLink(r3, s3)
69     net.addLink(s3, r4)
70     net.addLink(r4, s2)
71
72     # configure IPv6 addresses and forwarding table entries here
73
74     # set ipv6 addr and gateway
75     print "setup ipv6"
76     print set_ipv6(h1, "h1-eth0", h1_eth0_addr)
77     print set_ipv6(r1, "r1-eth0", r1_eth0_addr)
78     print set_ipv6(r1, "r1-eth1", r1_eth1_addr)
79     print set_ipv6(r2, "r2-eth0", r2_eth0_addr)
80     print set_ipv6(r2, "r2-eth1", r2_eth1_addr)
81     print set_ipv6(h2, "h2-eth0", h2_eth0_addr)
82     print set_ipv6(r3, "r3-eth0", r3_eth0_addr)
83     print set_ipv6(r3, "r3-eth1", r3_eth1_addr)
84     print set_ipv6(r4, "r4-eth0", r4_eth0_addr)
85     print set_ipv6(r4, "r4-eth1", r4_eth1_addr)
86     print "done setup ipv6"
87
88     # setup the local route
89     print "setup forward rules"
90     print set_ipv6_device_route(h1, subnet1_addr, "h1-eth0")
91     print set_ipv6_device_route(r1, subnet1_addr, "r1-eth0")
92     print set_ipv6_device_route(r1, subnet2_addr, "r1-eth1")
93     print set_ipv6_device_route(r2, subnet2_addr, "r2-eth0")
94     print set_ipv6_device_route(r2, subnet3_addr, "r2-eth1")
95     print set_ipv6_device_route(h2, subnet3_addr, "h2-eth0")
96     print set_ipv6_device_route(r3, subnet1_addr, "r3-eth0")
97     print set_ipv6_device_route(r3, subnet4_addr, "r3-eth1")
98     print set_ipv6_device_route(r4, subnet4_addr, "r4-eth0")
99     print set_ipv6_device_route(r4, subnet3_addr, "r4-eth1")
100    print "done setup forward rules"

```

```

101
102     # setup the gateway
103     print "setup gateway"
104     print set_ipv6_gateway(h1, r1_eth0_addr[:-3])
105     print set_ipv6_gateway(h2, r4_eth1_addr[:-3])
106     print "done setup gateway"
107
108     # setup router forward table
109     print "setup router forward table"
110     print set_ipv6_gateway_route(r1, subnet3_addr, r2_eth0_addr[:-3], "r1-eth1")
111     print set_ipv6_gateway_route(r2, subnet1_addr, r1_eth1_addr[:-3], "r2-eth0")
112     print set_ipv6_gateway_route(r3, subnet3_addr, r4_eth0_addr[:-3], "r3-eth1")
113     print set_ipv6_gateway_route(r4, subnet1_addr, r3_eth1_addr[:-3], "r4-eth0")
114     print "done setup router forward table"
115
116     # router forwarding
117     print "enable forwarding for routers"
118     r1.cmd("sysctl -w net.ipv6.conf.all.forwarding=1")
119     r2.cmd("sysctl -w net.ipv6.conf.all.forwarding=1")
120     r3.cmd("sysctl -w net.ipv6.conf.all.forwarding=1")
121     r4.cmd("sysctl -w net.ipv6.conf.all.forwarding=1")
122     print "done enabling forwarding for routers"
123
124     print h1.cmd("ip -V")
125
126
127     net.start()
128     CLI(net)
129     net.stop()

```

c)

We continue to run the mininet of Listing 4, and we manually set the mtu of the interface on r1(a.k.a r1-eth1) and r2(a.k.a r2-eth0), and using ping6 to ping h2 from h1 with a package size of 1450.(default MTU is 1500 on h1)

Listing 5: mininet result of c)

```

1 mininet> r1 ifconfig r1-eth1 mtu 1400
2 mininet> r2 ifconfig r2-eth0 mtu 1400
3 mininet> h1 ifconfig
4 h1-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
5         inet6 fe80::f8fb:27ff:fe5a:ef33 prefixlen 64 scopeid 0x20<link>
6         inet6 2001:638:709:a::1 prefixlen 64 scopeid 0x0<global>
7         ether fa:fb:27:5a:ef:33 txqueuelen 1000 (Ethernet)
8         RX packets 71 bytes 7026 (7.0 KB)
9         RX errors 0 dropped 0 overruns 0 frame 0
10        TX packets 45 bytes 4282 (4.2 KB)
11        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
12
13 lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
14        inet 127.0.0.1 netmask 255.0.0.0
15        inet6 ::1 prefixlen 128 scopeid 0x10<host>
16        loop txqueuelen 1000 (Local Loopback)
17        RX packets 12 bytes 1716 (1.7 KB)
18        RX errors 0 dropped 0 overruns 0 frame 0
19        TX packets 12 bytes 1716 (1.7 KB)
20        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

```

21 mininet> h2 ifconfig
22 h2-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
23     inet6 fe80::c041:eeff:fe30:8074 prefixlen 64 scopeid 0x20<link>
24     inet6 2001:638:709:b::1 prefixlen 64 scopeid 0x0<global>
25     ether c2:41:ee:30:80:74 txqueuelen 1000 (Ethernet)
26     RX packets 74 bytes 7284 (7.2 KB)
27     RX errors 0 dropped 0 overruns 0 frame 0
28     TX packets 47 bytes 4454 (4.4 KB)
29     TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
30
31 lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
32     inet 127.0.0.1 netmask 255.0.0.0
33     inet6 ::1 prefixlen 128 scopeid 0x10<host>
34     loop txqueuelen 1000 (Local Loopback)
35     RX packets 12 bytes 1716 (1.7 KB)
36     RX errors 0 dropped 0 overruns 0 frame 0
37     TX packets 12 bytes 1716 (1.7 KB)
38     TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
39
40 mininet> h1 tcpdump -i h1-eth0 -s 65535 -w ping1450mtu.dump &
41 tcpdump: h1 ping6 -M do -s 1450 2001:638:709:b::1
42 listening on h1-eth0, link-type EN10MB (Ethernet), capture size 65535 bytes
43 PING 2001:638:709:b::1(2001:638:709:b::1) 1450 data bytes
44 From 2001:638:709:a::f icmp_seq=1 Packet too big: mtu=1400
45 1458 bytes from 2001:638:709:b::1: icmp_seq=2 ttl=62 time=0.208 ms
46 1458 bytes from 2001:638:709:b::1: icmp_seq=3 ttl=62 time=0.094 ms
47 1458 bytes from 2001:638:709:b::1: icmp_seq=4 ttl=62 time=0.083 ms
48 1458 bytes from 2001:638:709:b::1: icmp_seq=5 ttl=62 time=0.085 ms
49 1458 bytes from 2001:638:709:b::1: icmp_seq=6 ttl=62 time=0.081 ms
50 1458 bytes from 2001:638:709:b::1: icmp_seq=7 ttl=62 time=0.084 ms
51 1458 bytes from 2001:638:709:b::1: icmp_seq=8 ttl=62 time=0.083 ms
52 ^C40 packets captured
53 40 packets received by filter
54 0 packets dropped by kernel

```

From the command line output, we can see that there is a message indicating that the package is too big.

We import the dump file 'ping1450mtu.dump' into the wireshark. We can clearly observe that all the ping6 packets are using ICMPv6 protocol. The first echo ICMPv6 packet has 1450 bytes of data and the first echo request didn't receive the echo response.

However, right after the first echo packet, h1(2001:638:709:a::1) receive a ICMPv6 packet from r1(2001:638:709:a::f) indicating that the packet is too big, with MTU of 1400.

Afterwards, the echo ICMPv6 packet is transferred in 2 IPv6 packets, and we can see that the ICMPv6 packet has been fragmented into 2 packets.

#### d)

We continue from previous problem, running the tracepath:

Listing 6: Running tracepath

```

1 mininet> h1 tracepath -n 2001:638:709:b::1
2 1?: [LOCALHOST] 1.176ms pmtu 1500
3 1: 2001:638:709:a::f 0.082ms
4 1: 2001:638:709:a::f 0.016ms
5 2: 2001:638:709:a::f 0.020ms pmtu 1400
6 2: 2001:638:709:f::2 0.064ms

```

```

7  3: 2001:638:709:b::1                                0.150ms reached
8  Resume: pmtu 1400 hops 3 back 3 1?: [LOCALHOST]      0.552ms pmtu 1500

```

---

Ping can determine the MTU value because if the current MTU is bigger, the packet will be dropped in the route and an ICMPv6 will be forwarded back including the error message and the supported MTU value.

Tracepath, just like traceroute[1], using the TTL and the error ICMPv6 packet to determine the right MTU value. It set different TTL value so that the packet only perform limited hops. Using this technique and in combination of the error packet sent back when MTU has exceeded and dropped, we can determine the link that limit the MTU and the optimal MTU value.

e)

No, we can do a experiment in the mininet

Listing 7: Path MTU cache experiment

---

```

1  mininet> h1 tracepath -n 2001:638:709:b::1
2  1?: [LOCALHOST]                                0.016ms pmtu 1500
3  1: 2001:638:709:a::f                            0.048ms
4  1: 2001:638:709:a::f                            0.015ms
5  2: 2001:638:709:a::f                            0.015ms pmtu 1400
6  2: 2001:638:709:f::2                            0.026ms
7  3: 2001:638:709:b::1                            0.037ms reached
8  Resume: pmtu 1400 hops 3 back 3
9  mininet> h1 tracepath -n 2001:638:709:b::1
10 1?: [LOCALHOST]                                0.016ms pmtu 1400
11 1: 2001:638:709:a::f                            0.042ms
12 1: 2001:638:709:a::f                            0.016ms
13 2: 2001:638:709:f::2                            0.025ms
14 3: 2001:638:709:b::1                            0.033ms reached
15 Resume: pmtu 1400 hops 3 back 3

```

---

As we can see, in the second tracepath, the pmtu value starts from 1400. Linux has obviously cached the path MTU. Such a cache is stored in the routing table entries. The host will not have routing table entries for every destination, but it can cache per-host route for every active destination[2].

## References

- [1] How does traceroute work and example's of using traceroute command.
- [2] Jeffrey Mogul and Steve Deering. Path mtu discovery. RFC 1191, RFC Editor, November 1990. <http://www.rfc-editor.org/rfc/rfc1191.txt>.